

# Jonathan Medhanie

604-440-5769 | johnnymedhanie@gmail.com

## EDUCATION

**UNIVERSITY OF  
BRITISH COLUMBIA**  
Bachelor's Degree,  
Computer Science &  
Mathematics  
Vancouver, BC  
Sept 2016 - April 2021

## SKILLS

C++ • Java • Python •  
Machine Learning •  
Numerical Methods • SQL  
• Docker • Golang •  
Distributed Systems •  
Hadoop •

## COURSEWORK

Algorithm Design (Graduate)  
Computational Optimization  
Computer Vision  
Geometric Modelling  
Relational Databases  
Software Engineering  
Applied Linear Algebra  
Complex Analysis  
Discrete Math (Honours)  
Graph Optimizations  
Multivariable Calculus  
Number Theory

## LINKS

Website:

[johnnymedhanie.github.io](http://johnnymedhanie.github.io)

Github:

[github.com/johnnymedhanie](https://github.com/johnnymedhanie)

Linkedin:

[in/jonathanmedhanie](https://in/jonathanmedhanie)

## WORK EXPERIENCE

**Software Development Engineer** | Amazon, *Alexa Skills*

August 2020 - Current

- Currently working on the Metrics analytics pipeline which allows insight into customer usage during Alexa skills development cycle

**Software Engineer Intern** | Twitter, *Real Time Storage*

June 2019 - August 2019

- Redesigned [Manhattan](#)'s data Exporter, Twitter's distributed database that stores all tweets and user data while serving 10 million queries per second across the world
- Optimized Exporter file count from **4.5M** to **20K** during Scan operation by implementing custom MapReduce jobs that splits and merges dataset partitions intelligently
- Tuned Hadoop file system and MapReduce configuration during MapReduce to improve CPU/memory use, reducing total job duration from **4 hours** to **20 minutes** on Twitter's largest cluster

**Software Engineer Intern** | Cisco Systems, *Cloud Infrastructure*

April 2018 - August 2018

- Optimized in-house AWS Virtual Machine healthchecker by reducing number requests made by pattern matching ip addresses/port, reducing time to completion from **30 minutes** to **25 seconds**
- Implemented a distributed tracing system for our API endpoint, allowing observability into durations of each user's unique request
- Refactored production databases' schema minimizing query duration

## PROJECTS

**Sentify** | Music Sentiment Analysis with Python/R

June 2018 - September 2018

- Utilizes natural language, text processing and machine learning to plot the variance of emotions in a songs given any musical artist
- Powered by making calls to Spotify and Genius' API detailed lyric and music library, taking an average of both the song's valence and energy in lyrics and tune.
- Plots processed data to dynamic web page deployed on Heroku

**Moment** | React Powered Productivity Browser Extension

May 2018 - January 2019

- Extension that periodically notifies user of personal achievements meant to motivate user to continue improving
- Architected, tested and deployed our backend which served authenticated requests
- Integrated with Chrome API to trigger database requests, periodic user view updates and timed notifications

## EXPERIENCE

**O11ycon - System Observability Conference** | Speaker  
August 2018 | San Francisco, CA

- Selected to speak and present the flaws of monitoring tools